

# Logistic Regression Project

December 9, 2017

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
/usr/local/lib/python2.7/dist-packages/pandas/core/computation/__init__.py:18: UserWarning: The
The minimum supported version is 2.4.6
```

```
ver=ver, min_ver=_MIN_NUMEXPR_VERSION), UserWarning)
```

```
In [2]: ad_data = pd.read_csv('advertising.csv')
```

```
In [3]: ad_data.head()
```

```
Out[3]:
```

	Daily Time Spent on Site	Age	Area Income	Daily Internet Usage	\
0	68.95	35	61833.90	256.09	
1	80.23	31	68441.85	193.77	
2	69.47	26	59785.94	236.50	
3	74.15	29	54806.18	245.89	
4	68.37	35	73889.99	225.58	

	Ad Topic Line	City	Male	Country	\
0	Cloned 5thgeneration orchestration	Wrightburgh	0	Tunisia	
1	Monitored national standardization	West Jodi	1	Nauru	
2	Organic bottom-line service-desk	Davidton	0	San Marino	
3	Triple-buffered reciprocal time-frame	West Terrifurt	1	Italy	
4	Robust logistical utilization	South Manuel	0	Iceland	

	Timestamp	Clicked on Ad
0	2016-03-27 00:53:11	0
1	2016-04-04 01:39:02	0
2	2016-03-13 20:35:42	0
3	2016-01-10 02:31:19	0
4	2016-06-03 03:36:18	0

```
In [4]: ad_data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 10 columns):
Daily Time Spent on Site    1000 non-null float64
Age                        1000 non-null int64
Area Income                1000 non-null float64
Daily Internet Usage       1000 non-null float64
Ad Topic Line              1000 non-null object
City                      1000 non-null object
Male                      1000 non-null int64
Country                   1000 non-null object
Timestamp                  1000 non-null object
Clicked on Ad              1000 non-null int64
dtypes: float64(3), int64(3), object(4)
memory usage: 78.2+ KB

```

```
In [5]: ad_data.describe()
```

```


Out[5]:      Daily Time Spent on Site      Age  Area Income  \
count      1000.000000  1000.000000    1000.000000
mean         65.000200    36.009000   55000.000080
std          15.853615     8.785562   13414.634022
min          32.600000    19.000000   13996.500000
25%          51.360000    29.000000   47031.802500
50%          68.215000    35.000000   57012.300000
75%          78.547500    42.000000   65470.635000
max          91.430000    61.000000   79484.800000

      Daily Internet Usage      Male  Clicked on Ad
count      1000.000000  1000.000000    1000.000000
mean         180.000100     0.481000     0.500000
std          43.902339     0.499889     0.500250
min         104.780000     0.000000     0.000000
25%         138.830000     0.000000     0.000000
50%         183.130000     0.000000     0.500000
75%         218.792500     1.000000     1.000000
max         269.960000     1.000000     1.000000

```


```
In [6]: ad_data['Age'].hist(bins =30)
```

```
Out[6]: <matplotlib.axes._subplots.AxesSubplot at 0x7f8fb56df410>
```




```
In [7]: sns.jointplot('Age', 'Area Income', ad_data)
```

```
Out[7]: <seaborn.axisgrid.JointGrid at 0x7f8fb56af210>
```




```
In [8]: sns.jointplot('Age', 'Daily Time Spent on Site', ad_data, kind= 'kde', color= 'red')
```

```
Out[8]: <seaborn.axisgrid.JointGrid at 0x7f8fb1cc7110>
```



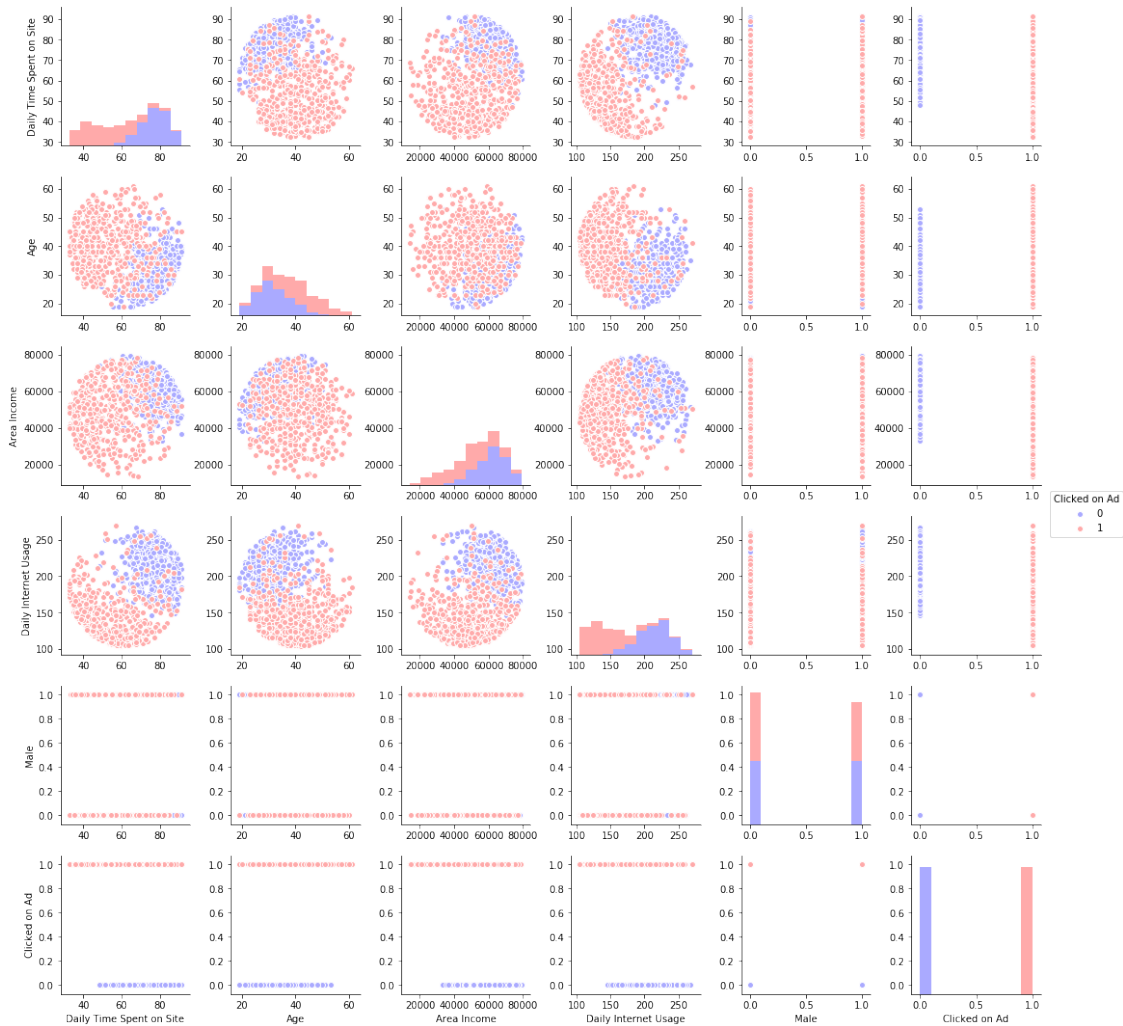
```
In [9]: sns.jointplot(x='Daily Time Spent on Site',y='Daily Internet Usage',data=ad_data,color='r')
```

```
Out[9]: <seaborn.axisgrid.JointGrid at 0x7f8fb56b8750>
```



```
In [10]: sns.pairplot(ad_data,hue='Clicked on Ad',palette='bwr')
```

```
Out[10]: <seaborn.axisgrid.PairGrid at 0x7f8fb18c8710>
```



```
In [11]: X = ad_data[['Daily Time Spent on Site', 'Age', 'Area Income', 'Daily Internet Usage', 'Male', 'Clicked on Ad']]
```

```
In [12]: y = ad_data['Clicked on Ad']
```

```
In [13]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
In [15]: from sklearn.linear_model import LogisticRegression
```

```
In [16]: model = LogisticRegression()
model.fit(X_train, y_train)
```

```
Out[16]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
verbose=0, warm_start=False)
```

```

In [17]: predictions= model.predict(X_test)

In [18]: from sklearn.metrics import classification_report

In [19]: print(classification_report(y_test,predictions))

```

	precision	recall	f1-score	support
0	0.84	0.97	0.90	146
1	0.96	0.82	0.89	154
avg / total	0.90	0.89	0.89	300